

Granularity versus Parallelization Overhead in High-Performance Computing Systems

Dongya Koh

University of Arkansas

October 9-10, 2018

MOTIVATION

- ▶ Solving a dynamic model with various kinds of heterogeneity needs to rely on numerical methods.
- ▶ Over the decades, a number of alternative solution methods were developed to improve the computation time and accuracy.
- ▶ The advent of multicore processors has accelerated the use of parallel computation.
- ▶ To further speed up, high-performance computing (HPC) system is now available.

A Basic Question:

- ▶ To speed up the computational time, how can we take advantage of HPC?

HOW CAN WE TAKE ADVANTAGE OF HPC?

The right choice of

1. Programming languages

C/C++, Fortran, Matlab, Python, Julia, R, etc.

2. Solution methods

VFI, PFI, Projection, Perturbation, EGM, etc.

3. Scalability (how effectively parallelizable)

4. Granularity (the amount of work in the parallel task)

5. Number of processing elements/cores

PROGRAMMING LANGUAGES

Language	Type	Rel. Exec. Time
C/C++	low-level, fastest with gcc	1.00
Fortran	low-level	1.05
Python	high-level, open-source, growing in popularity	$44\times \sim 270\times$
Julia	high-level, new open-source	$2.64\times \sim 2.70\times$
R	high-level, open-source	$281\times \sim 475\times$
Matlab	high-level, not free, license issue	$9\times \sim 11\times$
Mathematica	high-level	$809\times$

Source: Arouba & Fernandez-Villaverde (2014)

SOLUTION METHODS

Low Dimension	High Dimension
Value function iteration (VFI)	Smolyak sparse grid method (Krueger & Kubler, 2004, etc.)
Policy function iteration (PFI)	Adaptive grid method (Brumm & Scheidegger, 2017)
Projection method (Judd, 1992, etc.)	Stochastic simulation algorithm (Den Haan & Marcet, 1990, etc.)
Endogenous grid method (EGM) (Carroll, 2005, etc.)	ε -distinguishable set method (Judd, Maliar, Maliar, 2015)
Envelope condition method (Maliar & Maliar, 2013)	Cluster grid method (Judd, Maliar, Maliar, 2015)
Precomputation method (Judd, Maliar, Maliar, & Tsener, 2017)	Perturbation method (Judd & Guu, 1993, etc.)

OPTIMAL NUMBER OF (PHYSICAL OR VIRTUAL) CORES

- ▶ The total runtime of a program with n cores:

$$T(n) = \frac{T_p}{n} + T_s + P(n)$$

- ▶ The optimal number of cores to minimize the runtime:

$$P'(n^*) = \frac{T_p}{n^{*2}}$$

- ▶ If the parallelization overhead is approximated as $P(n) = pn^\alpha$ ($p, \alpha > 0$):

$$n^* = \left(\frac{T_p}{\alpha p} \right)^{1+\alpha}$$

PARALLELIZATION OVERHEAD

- ▶ Sources of parallelization overhead:
 1. Communication overhead in the form of synchronization and data communications
 2. Idling due to load imbalances
- ▶ Overheads vary by the implemented parallel algorithms and problems to be solved
- ▶ Parallelization overhead with n cores can be calculated as

$$P(n) = (T(n) - T_s) - \frac{T_p}{n},$$

provided that $P(1) = 0$

OVERVIEW

What We Do:

- ▶ Solve a variety of life-cycle models with efficient/inefficient solution methods.
- ▶ Compute the parallelization overhead.
- ▶ Compare the runtime of each parallel program by granularity and the number of cores.

What We Find:

- ▶ Parallelization overhead sets an upper bound to the HPC speed-ups.
- ▶ A design of HPC-efficient algorithm to reduce the parallelization overhead is essential for speed-ups in HPC.

HPC ENVIRONMENT

HPC System:

- ▶ HPC consists of three sub clusters
- ▶ Interconnected with a 324-port QDR 40 Gbps nonblocking QLogic Infiniband switch and supplementary switches
- ▶ 88TB long-term storage and 35 TB of scratch storage
- ▶ OS: Centos 6.5

Specifications of General Use Queue:

- ▶ CPU: 4x AMD Opteron 16-core 2.3 GHz 6276
- ▶ Memory/node: 512GB
- ▶ Max PBS spec: nodes=2:ppn=64
- ▶ Max PBS time: 72:00:00
- ▶ Software: Python/3.6.0-Anaconda

MODEL 1: CONSUMPTION/SAVING

- ▶ The baseline model assumes that an individual lives until age T .
- ▶ The individual's problem is to choose an amount of saving and consumption

$$V_t(a_t, e_t) = \max_{c_t, a_{t+1}} u(c_t) + \beta \mathbb{E}_t V_{t+1}(a_{t+1}, e_{t+1})$$

$$\text{s.t.} \quad c_t + a_{t+1} = we_t + (1+r)a_t$$

$$a_{t+1} \geq \underline{a}$$

$$e_{t+1} \sim P(e_{t+1}|e_t).$$

- ▶ The utility function takes isoelastic preference, $u(c) = \frac{c^{1-\sigma}}{1-\sigma}$

MODEL 1: VFI ALGORITHM

Step 1. Initialization

- a. Set model parameters.
- b. Define grid points for $(a_t, e_t) \in \mathcal{A} \otimes \mathcal{E}$.
- c. Construct a transition matrix $P(e_{t+1}|e_t)$.

Step 2. Computing a household problem at $t = T$

- a. For any given (a_T, e_T) , $a_{T+1} = 0$. Therefore, we obtain

$$c_T = we_T + (1 + r)a_T$$

$$V_T(a_T, e_T) = u(c_T)$$

Step 3. Computing a household problem at $t < T$

- a. For each point (a_t, e_t) with V_{t+1} , compute for all a_{t+1} ,
$$W_t(a_t, e_t) = u(we_t + (1 + r)a_t - a_{t+1}) + \beta \mathbb{E}_t V_{t+1}(a_{t+1}, e_{t+1})$$
 - b. Then, choose
$$\max_{a_{t+1}} W_t(a_t, e_t) = V_t(a_t, e_t).$$
-

MODEL 1: VFI ALGORITHM

Step 1. Initialization

- a. Set model parameters.
- b. Define grid points for $(a_t, e_t) \in \mathcal{A} \otimes \mathcal{E}$.
- c. Construct a transition matrix $P(e_{t+1}|e_t)$.

Step 2. Computing a household problem at $t = T$

- a. For any given (a_T, e_T) , $a_{T+1} = 0$. Therefore, we obtain

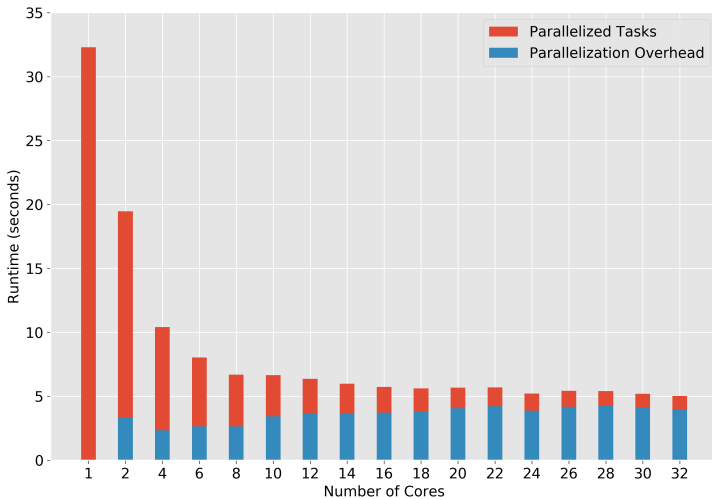
$$c_T = we_T + (1 + r)a_T$$

$$V_T(a_T, e_T) = u(c_T)$$

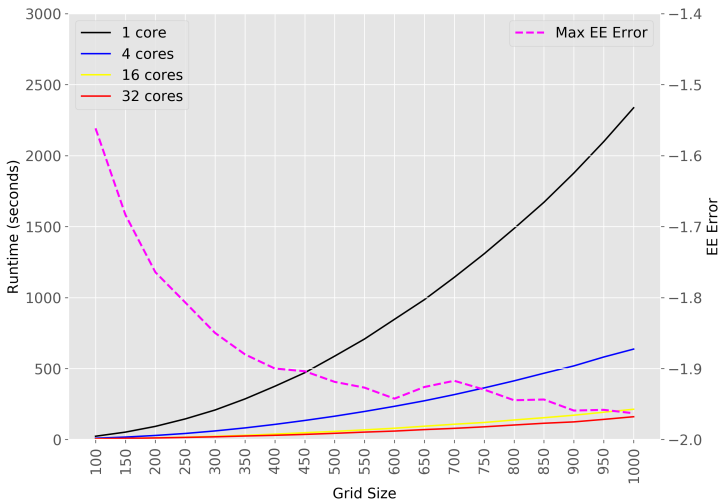
Step 3. Computing a household problem at $t < T$

- a. For each point (a_t, e_t) with V_{t+1} , compute for all a_{t+1} ,
$$W_t(a_t, e_t) = u(we_t + (1 + r)a_t - a_{t+1}) + \beta \mathbb{E}_t V_{t+1}(a_{t+1}, e_{t+1})$$
 - b. Then, choose
$$\max_{a_{t+1}} W_t(a_t, e_t) = V_t(a_t, e_t).$$
-

MODEL 1: VFI RUNNING TIME BY CORES



MODEL 1: VFI RUNNING TIME AND ACCURACY BY GRIDS



MODEL 1: EGM ALGORITHM

Step 1. Initialization

...

Step 2. (EXGM) Computing a household problem at $t = T$

- a. For any given (a_T, e_T) , $a_{T+1} = 0$. Therefore, we obtain

$$c_T = we_T + (1 + r)a_T$$

$$V_T^a = u'(c_T)(1 + r)$$

Step 3. (ENGM/EXGM) Computing a household problem at $t < T$

- a. (FOC) $\beta \mathbb{E} V_{t+1}^a = u'(c_t)$: compute \hat{c}_t for each point in (a_{t+1}, e_t) .
 b. (BC) $c_t = we_t + (1 + r)a_t - a_{t+1}$: compute \hat{a}_t for each $(\hat{c}_t, a_{t+1}, e_t)$.
 c. Given the obtained policy functions $g_t^c(\hat{a}_t, e_t) = \hat{c}_t$, $g_{t+1}^a(\hat{a}_t, e_t) = a_{t+1}$, interpolate new policy functions for each grid point (a_t, e_t) .
 d. (EC) $V_t^a = u'(c_t)(1 + r)$, we compute V_t^a .
-

MODEL 1: EGM ALGORITHM

Step 1. Initialization

...

Step 2. (EXGM) Computing a household problem at $t = T$

- a. For any given (a_T, e_T) , $a_{T+1} = 0$. Therefore, we obtain


$$c_T = we_T + (1 + r)a_T$$

$$V_T^a = u'(c_T)(1 + r)$$


Step 3. (ENGM/EXGM) Computing a household problem at $t < T$

- a. (FOC) $\beta \mathbb{E} V_{t+1}^a = u'(c_t)$: compute \hat{c}_t for each point in (a_{t+1}, e_t) .
 b. (BC) $c_t = we_t + (1 + r)a_t - a_{t+1}$: compute \hat{a}_t for each $(\hat{c}_t, a_{t+1}, e_t)$.
 c. Given the obtained policy functions $g_t^c(\hat{a}_t, e_t) = \hat{c}_t$, $g_{t+1}^a(\hat{a}_t, e_t) = a_{t+1}$, **interpolate new policy functions for each grid point (a_t, e_t) .**
 d. (EC) $V_t^a = u'(c_t)(1 + r)$, we compute V_t^a .
-

MODEL 1: EGM RUNNING TIME BY CORES



MODEL 1: EGM RUNNING TIME AND ACCURACY BY GRIDS



MODEL 2: ELASTIC LABOR SUPPLY

- The individual's problem is to choose an amount of saving, consumption, and labor supply.

$$V_t(a_t, e_t) = \max_{c_t, n_t, a_{t+1}} u(c_t, n_t) + \beta \mathbb{E}_t V_{t+1}(a_{t+1}, e_{t+1})$$

$$\text{s.t.} \quad c_t + a_{t+1} = we_t n_t + (1+r)a_t$$

$$a_{t+1} \geq \underline{a}$$

$$e_{t+1} \sim P(e_{t+1}|e_t).$$

- the utility function takes isoelastic preference,
$$u(c, n) = \frac{c^{1-\sigma}}{1-\sigma} - \chi \frac{n^{1+\eta}}{1+\eta}.$$

MODEL 2: EGM ALGORITHM

Step 1. Initialization

...

Step 2. (EXGM) Computing a household problem at $t = T$

- For each point (a_T, e_T) , $a_{T+1} = 0$.
- (FOC1) $\chi n_T^\eta = we_T(we_T n_T + (1+r)a_T)^{-\sigma}$: solve a nonlinear problem for \hat{n}_T
- (BC) $c_T = we_T n_T + (1+r)a_T$: compute \hat{c}_T .
- (EC) $V_T^a(a_T, e_T) = u_c(c_T, n_T)(1+r)$: compute V_T^a .

Step 3. (ENGM/EXGM) Computing a household problem at $t < T$

- (FOC2) $c_t^{-\sigma} = \beta \mathbb{E} V_{t+1}^a$: compute \hat{c}_t for each point in (a_{t+1}, e_t) .
- From (FOC1), compute \hat{n}_t .
- (BC) $a_{t+1} = we_t n_t + (1+r)a_t - c_t$: compute \hat{a}_t .
- Given the obtained policy functions $g_t^c(\hat{a}_t, e_t) = \hat{c}_t$, $g_t^n(\hat{a}_t, e_t) = \hat{n}_t$, $g_{t+1}^a(\hat{a}_t, e_t) = a_{t+1}$, we interpolate new policy functions, $\tilde{g}_t^c(a_t, e_t) = \hat{c}_t$, $\tilde{g}_t^n(a_t, e_t) = \hat{n}_t$, and $\tilde{g}_{t+1}^a(a_t, e_t) = \tilde{a}_{t+1}$ for each grid point $(a_t, e_t) \in \mathcal{A} \otimes \mathcal{E}$.
- If $\tilde{g}_{t+1}^a(a_t, e_t) = \tilde{a}_{t+1} < \underline{a}$, then do Step 2 only for the specific $(a_t, e_t) \in \mathcal{A} \otimes \mathcal{E}$. Otherwise, we compute V_t^a from (EC).

MODEL 2: EGM ALGORITHM

Step 1. Initialization

...

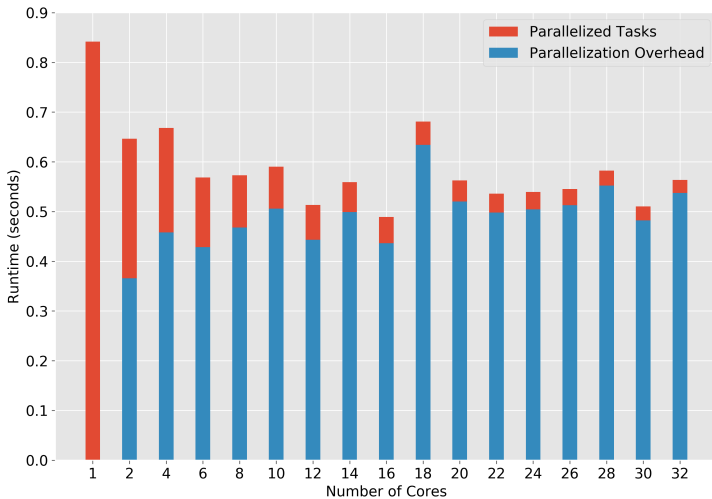
Step 2. (EXGM) Computing a household problem at $t = T$

- a. For each point (a_T, e_T) , $a_{T+1} = 0$.
- b. (FOC1) $\chi n_T^\eta = we_T (we_T n_T + (1+r)a_T)^{-\sigma}$: solve a nonlinear problem for \hat{n}_T
- c. (BC) $c_T = we_T n_T + (1+r)a_T$: compute \hat{c}_T .
- d. (EC) $V_T^a(a_T, e_T) = u_c(c_T, n_T)(1+r)$: compute V_T^a .


Step 3. (ENGM/EXGM) Computing a household problem at $t < T$

- a. (FOC2) $c_t^{-\sigma} = \beta \mathbb{E} V_{t+1}^a$: compute \hat{c}_t for each point in (a_{t+1}, e_t) .
- b. From (FOC1), compute \hat{n}_t .
- c. (BC) $a_{t+1} = we_t n_t + (1+r)a_t - c_t$: compute \hat{a}_t .
- d. Given the obtained policy functions $g_t^c(\hat{a}_t, e_t) = \hat{c}_t$, $g_t^n(\hat{a}_t, e_t) = \hat{n}_t$, $g_{t+1}^a(\hat{a}_t, e_t) = a_{t+1}$, we interpolate new policy functions, $\tilde{g}_t^c(a_t, e_t) = \hat{c}_t$, $\tilde{g}_t^n(a_t, e_t) = \hat{n}_t$, and $\tilde{g}_{t+1}^a(a_t, e_t) = \hat{a}_{t+1}$ for each grid point $(a_t, e_t) \in \mathcal{A} \otimes \mathcal{E}$.
- e. If $g_{t+1}^a(a_t, e_t) = \tilde{a}_{t+1} < \underline{a}$, then do Step 2 only for the specific $(a_t, e_t) \in \mathcal{A} \otimes \mathcal{E}$. Otherwise, we compute V_t^a from (EC).

MODEL 2: EGM RUNNING TIME BY CORES



MODEL 2: EGM RUNNING TIME AND ACCURACY BY GRIDS



MODEL 3: HUMAN CAPITAL INVESTMENT

- The individual's problem is to choose an amount of saving, consumption, and human capital investment

$$V_t(a_t, h_t, e_t) = \max_{c_t, a_{t+1}, s_t, h_{t+1}} u(c_t) + \beta \mathbb{E} V_{t+1}(a_{t+1}, h_{t+1}, e_{t+1})$$

$$\text{s.t.} \quad c_t + a_{t+1} = w e_t h_t (1 - s_t) + (1 + r) a_t$$

$$h_{t+1} = (1 - \delta) h_t + A_h h_t^\alpha s_t^\gamma$$

$$s_t \in [0, 1]$$

$$a_{t+1} \geq \underline{a}$$

$$e_{t+1} \sim P(e_{t+1} | e_t).$$

- The utility function takes isoelastic preference, $u(c) = \frac{c^{1-\sigma}}{1-\sigma}$

MODEL 3: EGM ALGORITHM

Step 1. Initialization

...

Step 2. (EXGM) Computing a household problem at $t = T$

- a. For each point (a_T, h_T, e_T) , $a_{T+1} = s_T = 0$.
 - b. (BC) $c_T = we_T h_T + (1 + r)a_T$: compute \hat{c}_T .
 - c. (EC1) $V_T^a = u'(c_T)(1 + r)$: compute V_T^a .
 - d. (EC2) $V_T^h = u'(c_T)we_T \left(1 - \left(1 - \frac{\alpha}{\gamma} \right) s_T + \frac{1 - \delta}{\gamma A_h h_T^{\alpha-1} s_T^{\gamma-1}} \right)$:
compute V_T^h .
-

MODEL 3: EGM ALGORITHM

Step 3. (ENGM/EXGM) Computing a household problem at $t < T$


- a. (FOC1) $\mathbb{E}V_{t+1}^a we_t h_t = \mathbb{E}V_{t+1}^h \gamma A_h h_t^\alpha s_t^{\gamma-1}$: solve a root-finding problem for \hat{s}_t for each point (a_{t+1}, h_t, e_t) .
 - b. (FOC2) $u'(c_t) = \beta \mathbb{E}V_{t+1}^a$: compute \hat{c}_t .
 - c. (BC) $c_t + a_{t+1} = we_t h_t (1 - s_t) + (1 + r)a_t$: compute \hat{a}_t .
 - d. Given the obtained policy functions $g_t^c(\hat{a}_t, h_t, e_t) = \hat{c}_t$, $g_t^s(\hat{a}_t, h_t, e_t) = \hat{s}_t$, $g_{t+1}^a(\hat{a}_t, h_t, e_t) = a_{t+1}$, interpolate new policy functions, $g_t^c(a_t, h_t, e_t) = \tilde{c}_t$, $g_t^s(a_t, h_t, e_t) = \tilde{s}_t$, and $g_{t+1}^a(a_t, h_t, e_t) = \tilde{a}_{t+1}$ for each grid point (a_t, h_t, e_t) .
 - e. If $g_{t+1}^a(a_t, h_t, e_t) = \tilde{a}_{t+1} < \underline{a}$, then do Step 3(a) only for the specific (a_t, h_t, e_t) and $a_{t+1} = \underline{a}$. Otherwise, we compute V_t^a and V_t^h from (EC1) and (EC2).
-

MODEL 3: EGM ALGORITHM

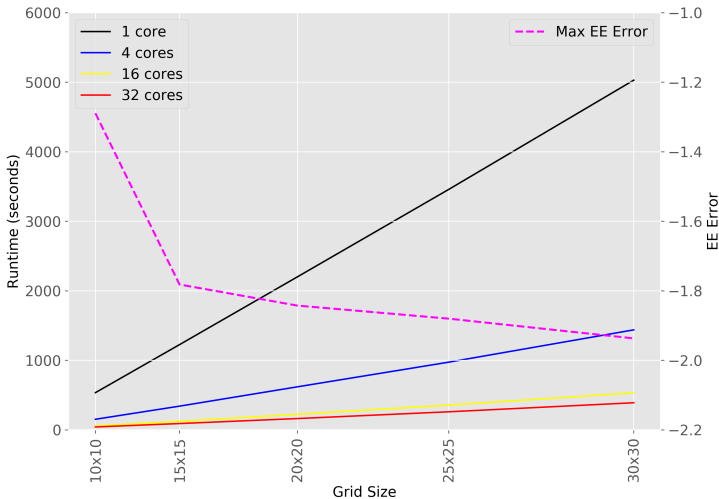
Step 3. (ENGM/EXGM) Computing a household problem at $t < T$

- a. (FOC1) $\mathbb{E}V_{t+1}^a we_t h_t = \mathbb{E}V_{t+1}^h \gamma A_h h_t^\alpha s_t^{\gamma-1}$: solve a root-finding problem for \hat{s}_t for each point (a_{t+1}, h_t, e_t) .
 - b. (FOC2) $u'(c_t) = \beta \mathbb{E}V_{t+1}^a$: compute \hat{c}_t .
 - c. (BC) $c_t + a_{t+1} = we_t h_t (1 - s_t) + (1 + r)a_t$: compute \hat{a}_t .
 - d. Given the obtained policy functions $g_t^c(\hat{a}_t, h_t, e_t) = \hat{c}_t$, $g_t^s(\hat{a}_t, h_t, e_t) = \hat{s}_t$, $g_{t+1}^a(\hat{a}_t, h_t, e_t) = a_{t+1}$, interpolate new policy functions, $\tilde{g}_t^c(a_t, h_t, e_t) = \hat{c}_t$, $\tilde{g}_t^s(a_t, h_t, e_t) = \hat{s}_t$, and $\tilde{g}_{t+1}^a(a_t, h_t, e_t) = \tilde{a}_{t+1}$ for each grid point (a_t, h_t, e_t) .
 - e. If $\tilde{g}_{t+1}^a(a_t, h_t, e_t) = \tilde{a}_{t+1} < \underline{a}$, then do Step 3(a) only for the specific (a_t, h_t, e_t) and $a_{t+1} = \underline{a}$. Otherwise, we compute V_t^a and V_t^h from (EC1) and (EC2).
-

MODEL 3: EGM RUNNING TIME BY CORES



MODEL 3: EGM RUNNING TIME AND ACCURACY BY GRIDS



TECHNICAL ISSUES USING MATLAB & PYTHON IN HPC

Matlab:

- ▶ The drawback of the Parallel Computing Toolbox is that it contains digital restrictions that only permit it to operate with 12 cores or fewer on a single computer.
- ▶ Matlab Distributed Computing Server costs several thousand dollars depending on the number of nodes licensed.

Python:

- ▶ A joint use of numba and multiprocessing packages has a systemic conflict.

CONCLUSION

- ▶ *Efficiently* running a parallel program in HPC environment requires a great amount of engineering effort.
- ▶ Parallelization overhead sets an upper bound to the HPC speed-ups.
- ▶ A design of HPC-efficient algorithm to reduce the parallelization overhead is essential for speed-ups in HPC.